

Machine Learning Course Project

January 27, 2016

Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement – a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://groupware.les.inf.puc-rio.br/har> (see the section on the Weight Lifting Exercise Dataset).

Data

The training data for this project are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

The test data are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

For ease of processing files and simplification of code, download the csv files and save to the set working directory.

```
setwd("~/Desktop/course/MachineLearning")
```

Load the libraries we'll be using:

```
## Load the preferred libraries
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.2.3
```

```
## Loading required package: lattice
## Loading required package: ggplot2
```

```
library(rpart)
library(plotmo)
```

```
## Loading required package: plotrix
```

```
## Warning: package 'plotrix' was built under R version 3.2.3
```

```
## Loading required package: TeachingDemos
```

```
library(randomForest)
```

```
## randomForest 4.6-12  
## Type rfNews() to see new features/changes/bug fixes.
```

```
library(AppliedPredictiveModeling)
```

With the following code, we'll read the data, and take a quick look at the properties. I also took a quick look at the .csv file by opening the test version in Excel.

```
pmlTrain <- read.csv("pml-training.csv", header=TRUE, na.strings=c("NA", "#Div/0!")) ## The training set  
pmlTest <- read.csv("pml-testing.csv", header=TRUE, na.string=c("NA", "#Div/0!")) ## The test set -set a  
dim(pmlTrain)
```

```
## [1] 19622 160
```

```
dim(pmlTest)
```

```
## [1] 20 160
```

```
summary(pmlTrain$classe)
```

```
##      A      B      C      D      E  
## 5580 3797 3422 3216 3607
```

In the training set, there are 19622 records with 159 variables (the first column is just a numeric count of observations). The “classe” variable we are solving for is divided among 5 classes.

The test set that was provided contains the exercise readings for 20 participants without the “classe” variable provided. We'll attempt to determine this “classe” with the use of predictive modelling, built using the training set. We'll set the test set to the side until the models are completed. At the end, we'll apply the same cleaning and transformations to that data, then apply our model.

Cleaning the data

We'll clean out the NearZeroVariance variables and remove them and the first column (a count of the observations) from our data as these will not contribute to the predictive model.

```
nzv <- nearZeroVar(pmlTrain, saveMetrics=TRUE) ## remove nearZeroVariances  
pmlTrain <- pmlTrain[,nzv$nzv==FALSE]  
pmlTrain <- pmlTrain[c(-1)] ## remove first column (count)
```

Remove observations with 75% NA:

```
noNAs<- pmlTrain ## find and remove 75% of NAs  
for(i in 1:length(pmlTrain)) {  
  if( sum( is.na( pmlTrain[, i] ) ) /nrow(pmlTrain) >= 0.75) {  
    for(j in 1:length(noNAs)) {  
      if( length( grep(names(pmlTrain[i]), names(noNAs)[j]) ) == 1) {
```

```

        noNAs <- noNAs[ , -j]
      }
    }
  }
}
pmlTrain <- noNAs ## set back to name
rm(noNAs) ## remove excess data
dim(pmlTrain)

```

```
## [1] 19622    58
```

This brings us down to 58 columns.

Split data

Now we'll split the data into a 60/40 training/test set to train the model then test the model before using for our prediction on the 20 observations in the final set.

```

inTrain <- createDataPartition(y=pmlTrain$classe,p=.60,list=FALSE)
train <-pmlTrain[inTrain,]
test <- pmlTrain[-inTrain,]

```

```
dim(train)
```

```
## [1] 11776    58
```

```
dim(test)
```

```
## [1] 7846    58
```

Prediction with Random Forests

For Random Forest information

```

set.seed(4726)
modelFitRF1 <- randomForest(classe ~ ., data=train)
predictionRF1 <- predict(modelFitRF1, test, type = "class")
modelRF <- confusionMatrix(predictionRF1, test$classe)
modelRF

```


```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 2232    1    0    0    0
##           B    0 1517    0    0    0
##           C    0    0 1365    2    0
##           D    0    0    3 1283    0
##           E    0    0    0    1 1442

```

```
##
## Overall Statistics
##
##           Accuracy : 0.9991
##           95% CI   : (0.9982, 0.9996)
##    No Information Rate : 0.2845
##    P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9989
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000  0.9993  0.9978  0.9977  1.0000
## Specificity      0.9998  1.0000  0.9997  0.9995  0.9998
## Pos Pred Value   0.9996  1.0000  0.9985  0.9977  0.9993
## Neg Pred Value   1.0000  0.9998  0.9995  0.9995  1.0000
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2845  0.1933  0.1740  0.1635  0.1838
## Detection Prevalence 0.2846  0.1933  0.1742  0.1639  0.1839
## Balanced Accuracy 0.9999  0.9997  0.9987  0.9986  0.9999
```

```
plot(modelFitRF1)
```



```
importance(modelFitRF1)
```

```
##           MeanDecreaseGini
```

## user_name	83.44910
## raw_timestamp_part_1	955.04706
## raw_timestamp_part_2	10.27330
## cvtd_timestamp	1398.27820
## num_window	593.69649
## roll_belt	541.82288
## pitch_belt	299.54879
## yaw_belt	351.81616
## total_accel_belt	112.56807
## gyros_belt_x	36.44526
## gyros_belt_y	53.04001
## gyros_belt_z	119.58138
## accel_belt_x	64.96234
## accel_belt_y	66.89162
## accel_belt_z	192.83989
## magnet_belt_x	105.66752
## magnet_belt_y	190.87321
## magnet_belt_z	166.96206
## roll_arm	121.17241
## pitch_arm	56.48048
## yaw_arm	73.75083
## total_accel_arm	31.03000
## gyros_arm_x	40.37999
## gyros_arm_y	45.21023
## gyros_arm_z	18.55010
## accel_arm_x	102.64520
## accel_arm_y	53.23422
## accel_arm_z	40.54220
## magnet_arm_x	89.34380
## magnet_arm_y	74.23590
## magnet_arm_z	58.02497
## roll_dumbbell	194.92801
## pitch_dumbbell	85.76263
## yaw_dumbbell	116.03882
## total_accel_dumbbell	124.97924
## gyros_dumbbell_x	39.60897
## gyros_dumbbell_y	90.25696
## gyros_dumbbell_z	24.61543
## accel_dumbbell_x	125.77462
## accel_dumbbell_y	170.73895
## accel_dumbbell_z	140.49083
## magnet_dumbbell_x	238.78228
## magnet_dumbbell_y	309.69705
## magnet_dumbbell_z	304.78809
## roll_forearm	219.52523
## pitch_forearm	305.94126
## yaw_forearm	50.67718
## total_accel_forearm	28.68905
## gyros_forearm_x	22.67795
## gyros_forearm_y	39.42047
## gyros_forearm_z	24.13554
## accel_forearm_x	133.13730
## accel_forearm_y	43.98957
## accel_forearm_z	91.79566

```
## magnet_forearm_x      81.03219
## magnet_forearm_y      66.71090
## magnet_forearm_z      88.31883
```

```
print(modelRF)
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 2232    1    0    0    0
##           B    0 1517    0    0    0
##           C    0    0 1365    2    0
##           D    0    0    3 1283    0
##           E    0    0    0    1 1442
##
## Overall Statistics
##
##           Accuracy : 0.9991
##           95% CI : (0.9982, 0.9996)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9989
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000  0.9993  0.9978  0.9977  1.0000
## Specificity      0.9998  1.0000  0.9997  0.9995  0.9998
## Pos Pred Value   0.9996  1.0000  0.9985  0.9977  0.9993
## Neg Pred Value   1.0000  0.9998  0.9995  0.9995  1.0000
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2845  0.1933  0.1740  0.1635  0.1838
## Detection Prevalence 0.2846  0.1933  0.1742  0.1639  0.1839
## Balanced Accuracy 0.9999  0.9997  0.9987  0.9986  0.9999
```

Prediction with Decision Trees

Create decision tree

Instead of plotting a decision tree, we can quickly look at a graph of the cross-validation results, and review the confusion matrix results and see that the error rate is higher than the random forest method.

```
set.seed(4726)
modelFitDT1 <- rpart(classe ~.,method="class", data=train)
predictionDT1 <- predict(modelFitDT1, test, type = "class")
modelDT <- confusionMatrix(predictionDT1,test$classe)
plotcp(modelFitDT1)
```



modelDT

Confusion Matrix and Statistics

##

Reference

Prediction	A	B	C	D	E
A	2151	53	10	1	0
B	54	1254	60	67	0
C	27	202	1272	206	63
D	0	9	26	959	181
E	0	0	0	53	1198

##

Overall Statistics

##

Accuracy : 0.871

95% CI : (0.8634, 0.8784)

No Information Rate : 0.2845

P-Value [Acc > NIR] : < 2.2e-16

##

Kappa : 0.837

McNemar's Test P-Value : NA

##

Statistics by Class:

##

	Class: A	Class: B	Class: C	Class: D	Class: E
Sensitivity	0.9637	0.8261	0.9298	0.7457	0.8308
Specificity	0.9886	0.9714	0.9231	0.9671	0.9917

## Pos Pred Value	0.9711	0.8739	0.7186	0.8162	0.9576
## Neg Pred Value	0.9856	0.9588	0.9842	0.9510	0.9630
## Prevalence	0.2845	0.1935	0.1744	0.1639	0.1838
## Detection Rate	0.2742	0.1598	0.1621	0.1222	0.1527
## Detection Prevalence	0.2823	0.1829	0.2256	0.1498	0.1594
## Balanced Accuracy	0.9762	0.8987	0.9265	0.8564	0.9113

A quick look at the results on the Decision Tree method and we see a lower accuracy rate, 88.73, than the Random Forest accuracy rate of 99.83% with a .17% for our out-of-sample error rate, so we'll progress with the Random Forest for our prediction set.

Predicting our results

First, we'll use the same cleaning methods as above:

```
cleanFormat <- colnames(train[,-58]) # classe column removal
pmlTest <- pmlTest[cleanFormat]
dim(pmlTest)
```

```
## [1] 20 57
```

And we'll coerce the data into the same format:

```
for (i in 1:length(pmlTest) ) {
  for(j in 1:length(train)) {
    if( length( grep(names(train[i]), names(pmlTest)[j]) ) == 1) {
      class(pmlTest[j]) <- class(train[i])
    }
  }
}

# To get the same class between pmlTest and train
pmlTest <- rbind(train[2,-58], pmlTest) ## remove excess rows
pmlTest <- pmlTest[-1,]
```

Then we apply the prediction model to the data:

```
predictionFinal <- predict(modelFitRF1, pmlTest, type="class")
```

And our final results for our 20 test cases.

```
predictionFinal
```

```
## 22 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21
## B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```